

REPORT

Aim

To understand the difference between how virtual memory is allocated and mapped using `kmalloc` and `vmalloc`.

Methodology

Software Walk

1. Given page VA, we can perform a software page table walk doing the following in order,
 - `pgd_offset()` - returns pgd offset given `mm_struct` of task and VA
 - `pud_offset()` - returns pud offset given VA and pgd offset
 - `pmd_offset()` - returns pmd offset given VA and pud offset
 - `pte_offset_map()` - returns pte given VA and pmd offset
2. If any of the above returns a none and bad value we exit with return value as 0
3. If everything executes successfully we obtain the struct page using `pte_page()` and using the page struct we obtain the Physical address of page (PA) using `page_to_phys()`
4. Now we can obtain the PA of VA by taking the last 12 bits of VA and Oring it with the page PA.
5. Now this PA is returned

Module (a) logic

1. Allocate memory for a user defined struct using `kmalloc()`/`vmalloc()`
2. Initialize the struct data to a test value
3. Type cast the pointer to the newly created struct to unsigned long which gives us the VA
4. Now display the VA
5. get the PA using the `virt_to_phys()` macro by passing pointer to struct and display it
6. get the PA using the software walk method as described above using VA

Module (b) logic

1. Keeping module above module loaded
2. Load a new module with above displayed VA as a parameter

3. Access the struct data that was input in the above module and print it
4. Now try to access the VA using the same above two methods
5. for `virt_to_phys()` make sure to cast the address it to struct type before invoking the macro

Experiments

1- Kmalloc

```
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo insmod part2_a.ko choice=0
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo insmod part2_b.ko va=18446612135152665992
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo rmmod part2_b.ko
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo rmmod part2_a.ko
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$
```

Illustration 1: commands to run module with `kmalloc()`

1. Insert module(a) with choice=0 (invoked `kmalloc`)
2. With keeping the above module running insert module(b)
3. It is observed that the output using both the `virt_to_phys()` and software walk produce the same correct output in both modules
4. The **PA obtained in all 4 cases are the same** and correct
5. The second module is rightly able to access the first module's struct data
6. It is to be noted that along the path of traversal of the software page table walk, the **pgd-offset value is different** (can be seen below) **for the two modules** but map to the same pud-offset and hence the walks there after are the same.

```
Feb 16 23:59:33 prashanth-VirtualBox kernel: [22523.665023] Part-2: -----INSERTED-----
Feb 16 23:59:33 prashanth-VirtualBox kernel: [22523.665028] Node-P: ffff8800a92f4588 VA: 18446612135152665992
Feb 16 23:59:33 prashanth-VirtualBox kernel: [22523.665030] pgd-off:ffff88008c9880 pud-off:ffff880001ff8010 pmd-off:ffff88015fffea48 pte:80000000a92f4163
Feb 16 23:59:33 prashanth-VirtualBox kernel: [22523.665031] Physical Address Obtained from virt_to_phys : 00000000a92f4588
Feb 16 23:59:33 prashanth-VirtualBox kernel: [22523.665032] Physical Address Obtained from software walk: 00000000a92f4588
Feb 16 23:59:54 prashanth-VirtualBox kernel: [22544.506937] Part-2-Access: -----INSERTED-----
Feb 16 23:59:54 prashanth-VirtualBox kernel: [22544.506943] Node-P: ffff8800a92f4588 VA: 18446612135152665992
Feb 16 23:59:54 prashanth-VirtualBox kernel: [22544.506945] Value at Node: 10
Feb 16 23:59:54 prashanth-VirtualBox kernel: [22544.506947] pgd-off:ffff8800a5084880 pud-off:ffff880001ff8010 pmd-off:ffff88015fffea48 pte:80000000a92f4163
Feb 16 23:59:54 prashanth-VirtualBox kernel: [22544.506949] Physical Address Obtained from virt_to_phys : 00000000a92f4588
Feb 16 23:59:54 prashanth-VirtualBox kernel: [22544.506951] Physical Address Obtained from software walk: 00000000a92f4588
Feb 17 00:00:03 prashanth-VirtualBox kernel: [22553.287929] Part-2-Access: -----REMOVED-----
Feb 17 00:00:05 prashanth-VirtualBox kernel: [22555.707895] Part-2: -----REMOVED-----
```

Illustration 2: output for `kmalloc()`

2- Vmalloc

```
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo insmod part2_a.ko choice=1
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo insmod part2_b.ko va=18446683600585457664
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo rmmod part2_b.ko
prashanth@prashanth-VirtualBox:~/Modules/a1/part2$ sudo rmmod part2_a.ko
```

Illustration 3: Commands to run modules with `vmalloc()`

1. Insert module(a) with choice=1 (invoked `vmalloc`)

2. With keeping the above module running insert module(b)
3. It is observed that the **only software walk fetches the same correct PA** in both modules, and **virt_to_phys()** fetches the wrong PA.
4. The second module is rightly able to access the first module's struct data with input VA
5. It is to be noted that along the path of traversal of the software page table walk, the **pgd-offset value is different** (can be seen below) **for the two modules** but map to the same pud-offset and hence the walks there after are the same.

```
Feb 16 23:57:52 prashanth-VirtualBox kernel: [22422.057947] Part-2: -----INSERTED-----
Feb 16 23:57:52 prashanth-VirtualBox kernel: [22422.057963] Node-P: fffffc90000eb8000 VA: 18446683600585457664
Feb 16 23:57:52 prashanth-VirtualBox kernel: [22422.057967] pgd-off:ffff8800a9340c90 pud-off:ffff88015a09a000 pmd-off:ffff88015a09b038 pte:800000014e745163
Feb 16 23:57:52 prashanth-VirtualBox kernel: [22422.057970] Physical Address Obtained from virt_to_phys : 0000410000eb8000
Feb 16 23:57:52 prashanth-VirtualBox kernel: [22422.057972] Physical Address Obtained from software walk: 000000014e745000
Feb 16 23:58:23 prashanth-VirtualBox kernel: [22452.990737] Part-2-Access: -----INSERTED-----
Feb 16 23:58:23 prashanth-VirtualBox kernel: [22452.990746] Node-P: fffffc90000eb8000 VA: 18446683600585457664
Feb 16 23:58:23 prashanth-VirtualBox kernel: [22452.990749] Value at Node: 10
Feb 16 23:58:23 prashanth-VirtualBox kernel: [22452.990754] pgd-off:ffff8800bc9c1c90 pud-off:ffff88015a09a000 pmd-off:ffff88015a09b038 pte:800000014e745163
Feb 16 23:58:23 prashanth-VirtualBox kernel: [22452.990757] Physical Address Obtained from virt_to_phys : 0000410000eb8000
Feb 16 23:58:23 prashanth-VirtualBox kernel: [22452.990759] Physical Address Obtained from software walk: 000000014e745000
Feb 16 23:58:28 prashanth-VirtualBox kernel: [22458.045997] Part-2-Access: -----REMOVED-----
Feb 16 23:58:33 prashanth-VirtualBox kernel: [22463.118761] Part-2: -----REMOVED-----
```

Illustration 4: Output for vmalloc()

Observations

1. While performing **software walks** on the kernel memory using different modules using kmalloc()/vmalloc() **always points us to the right PA**
2. The **pgd-offset** which performing software walks for the same kernel VA from different modules would be **different** but they ultimately map to the same pud-offset and there by inheritantly to the right PA

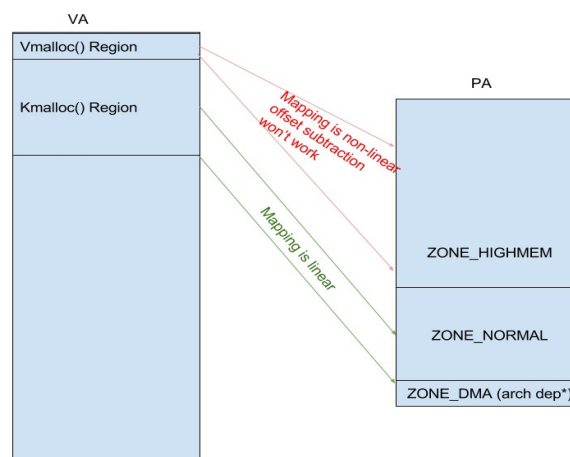


Illustration 5: Mapping of kmalloc vs vmalloc

3. **virt_to_phys()** macro fetches the **right PA while using kmalloc()** as it just subtracts the kernel memory offset for current architecture from VA to give its PA. And this will work for kmalloc as it is mapped on the zone_normal region of the physical address space with a linear offset.
4. **virt_to_phys()** macro fetches the **incorrect PA while using vmalloc()** as mapping is non-linear using vmalloc() is it may allocate space from any free space in highmem.